

Kunal Maheshwari

Haldwani, Uttarakhand, India | 9258446350 | work.kunal.maheshwari@gmail.com

LinkedIn: linkedin.com/in/kunalmaheshwari26 | GitHub: github.com/preferablehuman | Portfolio: preferablehuman.github.io/portfolio/

Summary

Data Engineering and Backend Developer with 4.6+ years of Java, SQL, cloud integration, CI/CD, and production-system experience, plus graduate AWS/Spark data-platform projects. Built event-driven pipelines using S3, Lambda, EFS, ECS, Spark Structured Streaming, RDS PostgreSQL, Terraform, and Docker; improved Spark processing 40%, optimized Oracle workflows 90%, and supported 50K+ daily API transactions at 99.99% SLA.

Work Experience

NTT Data, Bengaluru, India

Software Developer, Aug 2022 - Aug 2024

- Built healthcare integration services using Java 8/11, Spring Boot/Spring Framework, REST/SOAP APIs, SOA, IBM WebSphere, CAQH workflows, and API Gateway authentication for 50K+ daily requests at 99.99% SLA.
- Designed event-driven data propagation using Azure Relay, Azure Service Bus Topics, Azure Functions, and JMS queues across DEV, SIT, and PROD, improving reliability of downstream enterprise database workflows.
- Delivered Maven, Jenkins, Azure DevOps, UrbanCode, Git/SVN, JUnit, Log4j 2.x, and CI/CD release support, reducing issue-resolution time 20% and accelerating EDX platform delivery 50%.

Capgemini, Mumbai, India

Senior Analyst, Jul 2020 - Aug 2022

- Developed automotive Java/Spring backend services with REST APIs, JDBC, Oracle SQL, EJB, JSF, and high-availability transactional workflows, sustaining 99%+ uptime.
- Tuned Oracle SQL queries, database access paths, and legacy Java/J2EE components, improving production performance approximately 90% and strengthening data-backed service reliability.
- Integrated a Python decision-tree service through REST APIs for option availability, pricing, and business-rule validation while supporting Jenkins, GitLab CI/CD, Maven, JUnit4, and quality gates.

Selected Data Projects

NYU Taxi Data Analysis Platform

AWS, Spark, Terraform, PostgreSQL

- Built event-driven data pipeline where S3 uploads trigger Lambda, stage files into EFS, process micro-batches with Spark Structured Streaming on ECS, and write analytics to RDS PostgreSQL with checkpoints.
- Defined AWS infrastructure with Terraform across VPC, IAM, S3 notifications, Lambda, EFS access points, ECS services, RDS, and CloudWatch; processed 8GB+ taxi data with 40% runtime improvement.

Lamprey Ecological Tracking Pipeline

Python, Flask, OpenCV, Cloud SQL

- Processed video and tag-detection data using OpenCV, Pandas, Google Cloud Storage, Cloud SQL, PostgreSQL, SQLAlchemy, and Docker for academic ecological tracking workflows.

Technical Skills

Languages/Data: Python, SQL, Java 8/11, Pandas, NumPy, Jupyter.

Data Engineering: Apache Spark, PySpark, Spark Structured Streaming, ETL, Micro-batches, EFS Checkpoints, JDBC, PostgreSQL, Oracle SQL, RDS, Cloud SQL.

Cloud/DevOps: AWS S3, Lambda, EFS, ECS, RDS, CloudWatch, IAM, VPC, Terraform, Docker, Azure Service Bus, Azure Relay, Azure Functions, Jenkins, GitLab CI/CD, Maven, Git.

Backend/Quality: Spring Boot, REST/SOAP APIs, JMS, FastAPI, Flask, JUnit, Log4j 2.x, Agile/Scrum, Production Support, Performance Tuning.

Education and Certifications

Central Michigan University, MS Computer Science, Aug 2024 - Dec 2025, GPA: 3.87/4.0 | **Lovely Professional University**, B.Tech CSE, 2016 - 2020, CGPA: 8.02/10

Microsoft Certified: Azure Fundamentals | Azure Developer Associate | Azure Data Engineer Associate